

Apache Pig

Extract Transform Load

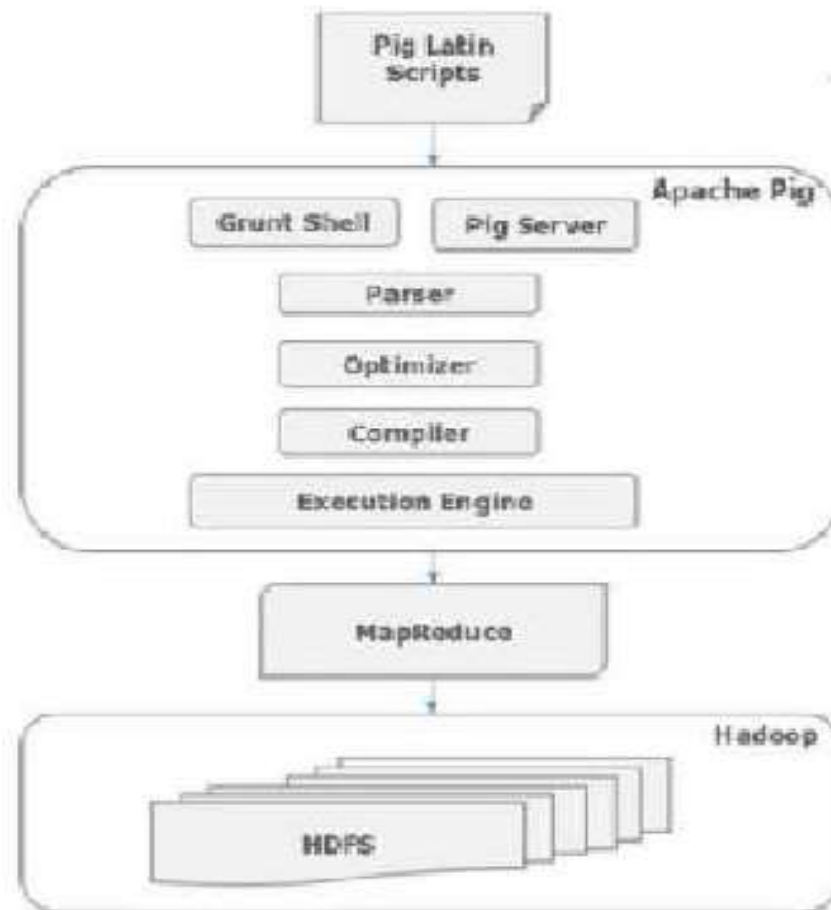
Dineshkumar S
SmartR.in

Introduction

Apache Pig:

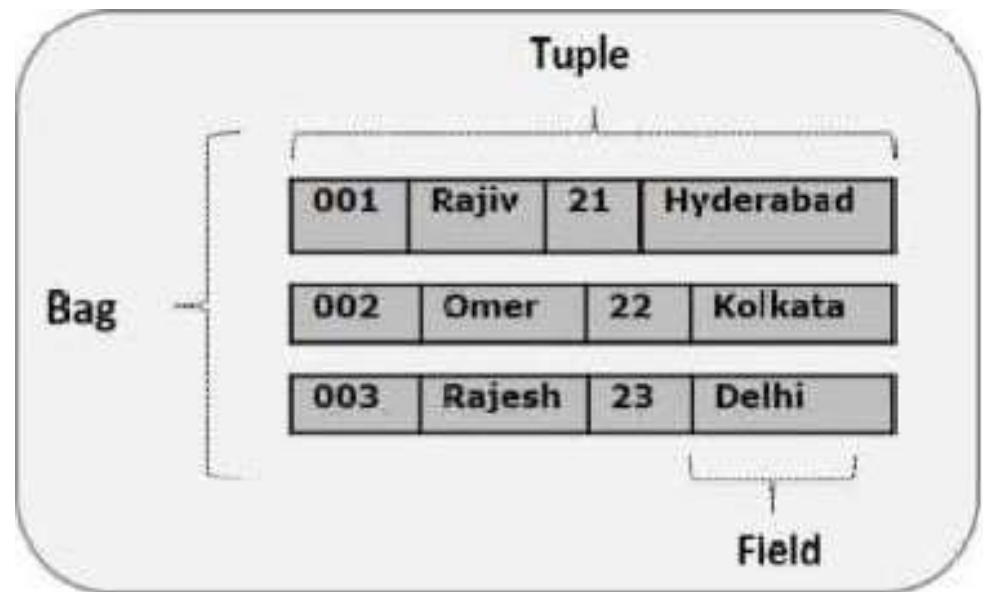
- is an abstraction over MapReduce
- has structure which is amenable for substantial parallelization
- is generally used with Hadoop
- does lazy evaluation
- uses Pig Latin, which provides
 - Ease of Programming
 - Optimization opportunities
 - Focus on semantics rather than efficiency
 - Extensibility
 - Can create UDFs

Architecture



Pig Latin Data Model

- Fully nested & complex data model
- Atom
- Tuple
- Bag
- Map
- Relation




Pig Environment

Establishing Pig environment contains following three major steps:

- Installation
 - <https://pig.apache.org/docs/r0.7.0/setup.html>
- Grunt shell
 - Pig interactive shell which accepts Pig-latin based statements.
- Execution
 - Two modes of execution, Mapreduce mode and Local mode, we'll see both ahead.

Pig Latin



Data Types

Basic	Complex
int, long, biginteger	Tuple : (raju, 5)
float, double, bigdecimal	Bag: {(raju,5),(ranu,6)}
chararray	Map: ['name'#'raju','age':40]
bytearray	
boolean	
datetime	

Operators and Constructs

- All standard Arithmetic and Relational operator
- Constructs
 - Tuple: ()
 - Bag: {}
 - Map: []
- Relational Operators
 - Load & Storing: **LOAD & STORE**
 - Filtering: **FILTER, DISTINCT, FOREACH, GENERATE, STREAM**
 - Grouping & Joining: **JOIN, COGROUP, CROSS, GROUP**
 - Sorting: **ORDER, LIMIT**
 - Combining & Splitting: **UNION, SPLIT**
 - Diagnostic & Debugging: **DUMP, DESCRIBE, EXPLAIN, ILLUSTRATE**

Sample Queries

```
student = LOAD 'hdfs://localhost:9000/pig_data/student_data.txt'  
  USING PigStorage(',')  
  as ( id:int, firstname:chararray, lastname:chararray, phone:chararray,  
    city:chararray );  
  
STORE student INTO ' hdfs://localhost:9000/pig_Output/' USING PigStorage (',');  
  
Dump student  
  
describe student;  
  
explain Relation_name;  
  
illustrate student  
  
group_data = GROUP student_details by age  
  
cogroup_data = COGROUP student_details by age, employee_details by age;  
  
result = JOIN relation1 BY columnname, relation2 BY columnname;
```

Sample Queries (contd.)

cross_data = CROSS customers, orders;

student = UNION student1, student2;

SPLIT student_details into student_details1 if age<23, student_details2 if (22<age and age>25);

filter_data = FILTER student_details BY city == 'Chennai';

distinct_data = DISTINCT student_details;

foreach_data = FOREACH student_details GENERATE id,age,city;

order_by_data = ORDER student_details BY age DESC;

limit_data = LIMIT student_details 4;

Eval & Load/Store Functions

Eval:

AVG()	BagToString()	CONCAT()	COUNT()	COUNT_STAR()	
DIFF()	IsEmpty()	MAX()	MIN()	SIZE()	SUM()

Load/Store:

PigStorage()

TextLoader()

BinStorage()

UDFs(User Defined Functions)

- UDF ?
 - Piggybanks?
 - Types of UDFs:
 - Filter Functions
 - Eval Functions
 - Algebraic Functions
 - { Writing UDFs
 - { Registering UDFs
 - { Defining Alias
 - { Using UDFs
- [documentation](#)

UDF's (contd.)

What is UDF?

In addition to the built-in functions, Pig provides excellent support for User Defined Functions(UDF's), where we can define our own functions and use them along with different Pig Latin built-in functions. Complete support to write UDF's is provided in Java and partially in other languages, viz. Jython, Python, JS, Ruby and Groovy.

Piggybanks?

Apache Pig provides a central repository for Java based UDF's using which you can use the UDF's written by other people and you can also contribute there.

UDF Example

- This example is written for converting given string to Uppercase,
- You can download the java code from [link](#).
- Extract it and create the jar file from the source code or download from [here](#).
- Now assume we have the jar file at /home/hduser/hadoop/exampleJARs/udfEx1.jar and some sample data at /hduser/sampleData/sampledData.txt
- Now on grunt shell

```
REGISTER /home/hduser/hadoop/exampleJARs/udfEx1.jar;
```

```
DEFINE udf_eval SampleEval();
```

```
data = load '/hduser/sampleData/sampledData.txt' using PigStorage(',') as (id:int, name:chararray, age:int, city:chararray);
```

```
upper_case_data = foreach data generate udf_eval(name);
```

```
store upper_case_data into 'hdfs://localhost:9000/hduser/outputudf15';
```

Pig Scripts

- Modes of running Pig Scripts
 - Mapreduce mode
 - Local mode (-x parameter)

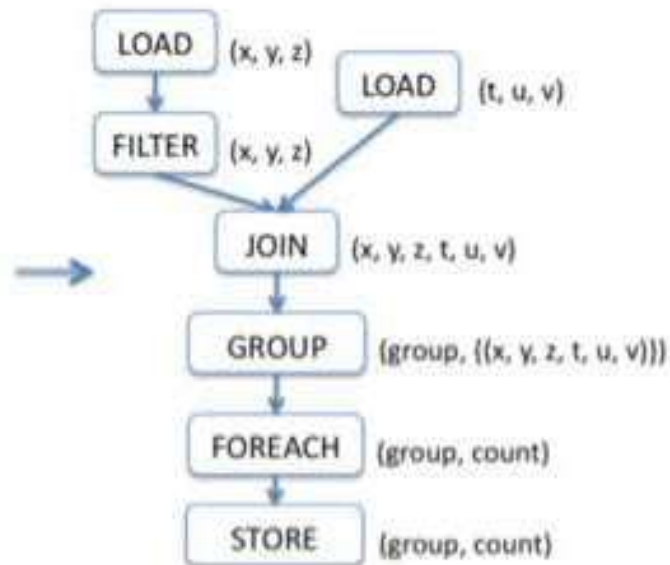
```
pig -x <mode> <pig_script_path>
```

Pig → MapReduce

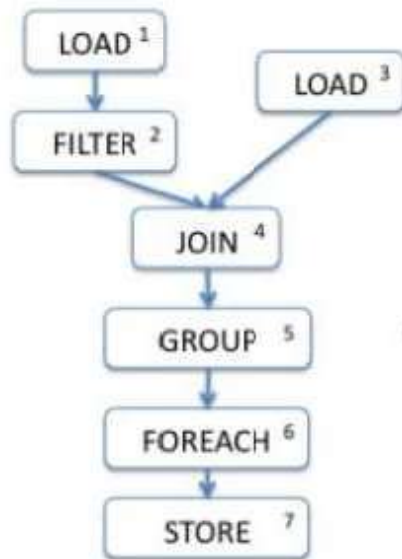
Pig Latin

```
A = LOAD 'file1' AS (x, y, z);  
B = LOAD 'file2' AS (t, u, v);  
C = FILTER A by y > 0;  
D = JOIN C BY x, B BY u;  
E = GROUP D BY z;  
F = FOREACH E GENERATE  
    group, COUNT(D);  
STORE F INTO 'output';
```

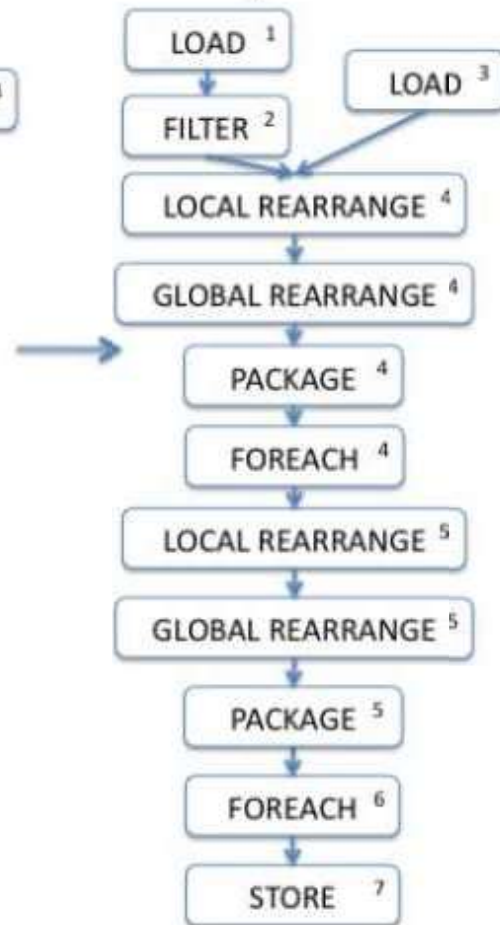
Logical Plan



Logical Plan



Physical Plan



Map Reduce Plan

